

CROP MONITORING DRONE

Simulation Validation Plan

A whitepaper-oriented experimental framework for quantifying the operational and economic value of RGB/NIR crop monitoring

Central hypothesis

High-frequency RGB/NIR UAV observations can reduce uncertainty in modeled crop state compared with conventional spatial sampling, enabling earlier anomaly detection and better crop-state/yield forecasts. When those observations trigger targeted scouting or management decisions, the added information can produce measurable labor, input, yield, and profit value.

*Prepared for the Locust Research crop-monitoring validation project
September 2026*

1. Executive Summary

The recommended validation framework treats the crop-monitoring drone as an information system rather than as an imaging product. The purpose of the simulation is to quantify how better spatial and temporal observations affect crop-state estimation, anomaly detection, management decisions, and ultimately farm economics.

Core methodological principle

Do not use the same crop model as both the simulated “truth” and the model being improved by drone observations in the final validation. Separate ground truth, imperfect observations, crop-state estimation, and decision logic so the experiment does not become circular.

The simulation should be built in three layers:

1. A spatial farm “truth” model that evolves crop, soil, weather, and stress conditions through time.
2. Observation models that simulate what a conventional scout and the RGB/NIR drone would actually perceive, including imperfect coverage and measurement error.
3. A decision/economic layer that converts observations into crop-model updates, scouting or treatment actions, yield outcomes, costs, and return on investment.

Recommended progression

Version	Primary question	Primary outputs
MVP / Version 1	Does drone-derived information improve crop-state and yield estimation versus conventional scouting?	LAI RMSE, yield RMSE, anomaly detection, time-to-detection, acres/hour, labor hours
Version 2	Does better monitoring lead to better management decisions?	Intervention timing, input use, avoided loss, net profit, ROI
Version 3	Can repeated UAV observations support predictive analytics and a farm “digital twin”?	Forecast accuracy, uncertainty reduction, risk prediction, optimal monitoring frequency


2. Simulation Architecture

The experiment should resemble an Observing System Simulation Experiment (OSSE): the true crop state exists inside the simulation but is hidden from both monitoring strategies. Each strategy receives only the data it could realistically obtain.

TRUE FARM STATE

(crop + soil + weather + management + stress)

|
v



2.1 Separation of truth and estimator

During development, WOFOST can be used to generate synthetic truth because it is convenient and allows the pipeline to be debugged. For the final whitepaper experiment, however, introduce model mismatch using one or more of the following:

- Generate the “true” crop with APSIM, DSSAT, or another process-based model while using WOFOST/PCSE as the estimator.
- Use substantially perturbed WOFOST parameters for truth versus estimation.
- Calibrate the truth layer against field data, then intentionally give the estimator imperfect soil, crop, or management parameters.
- Use historical weather years and heterogeneous soils not perfectly known by the estimator.

3. Phase 1 - Build the Spatial Ground-Truth Farm

Start with a single well-supported crop such as maize/corn. Use PCSE with WOFOST 8.1 where practical, particularly because the model includes crop growth, water balance, and nitrogen-related dynamics that are relevant to remote sensing and precision agriculture.

3.1 Spatial representation

Do not represent the entire farm with one crop-model run. Divide the field into grid cells or management zones. A practical MVP is 20-100 zones; a later implementation can use 5-20 m raster cells.

Variable	How it varies spatially	Why it matters
Soil texture / water holding capacity	By zone	Creates realistic water-stress patterns

Variable	How it varies spatially	Why it matters
Initial soil moisture	By zone and season	Changes crop response to dry periods
Nitrogen availability	By zone	Creates spectral/crop-growth differences
Plant population / emergence	By zone	Supports RGB stand-count and gap-detection cases
Planting date	Small offsets	Introduces phenology variability
Drainage / low areas	By zone	Supports waterlogging or delayed-growth scenarios

3.2 Core state variables

State	Role in simulation	Expected observability
Leaf Area Index (LAI)	Primary crop-state variable for assimilation	Estimated from RGB/NIR features
Biomass	Crop growth / yield pathway	Indirectly inferred
Development stage	Phenology and timing	Partly inferable; known from model/planting date
Soil moisture	Water-stress driver	Not directly measured by RGB/NIR
Crop nitrogen status	N-stress driver	Indirect spectral/canopy signal
Canopy cover / stand density	Structural crop condition	Strong RGB use case
Yield / storage organ weight	Primary end-of-season result	Forecast target, not directly observed

4. Phase 2 - Introduce Controlled Crop Problems

Create stress events with known onset, location, severity, duration, and true crop impact. This gives the experiment objective ground truth for detection probability and detection latency.

Stress case	Simulation suitability	Primary drone signal / value
Water stress	High	Growth divergence; vegetation response; anomaly mapping
Nitrogen stress	High with suitable WOFOST configuration	Changes in vigor, canopy and spectral response
Poor emergence / stand loss	High with external spatial layer	Direct RGB plant/row/gap detection
Generic canopy damage	High as external disturbance	Tests anomaly detection without overclaiming diagnosis
Disease / pest pressure	Later phase	Best modeled first as spatial canopy-health effect unless a disease model is available

Example event sequence:

- Day 35: Zone 18 begins experiencing water stress.
- Day 38: simulated LAI diverges from the unstressed trajectory.
- Day 40: the virtual RGB/NIR observation becomes measurably anomalous.
- Day 43: visible symptoms become clear to an observer.
- Day 48: a conventional scout samples the affected location.

Key KPI

Time-to-detection = detection date - stress onset date. This converts monitoring frequency and spatial coverage into an agronomically meaningful performance metric.

5. Phase 3 - Build a Virtual Version of the Actual Drone

The simulated platform should use measured engineering data from the real drone rather than generic industry assumptions. This is the bridge between physical testing and simulated farm value.

Engineering input	Use in simulation
Cruise speed	Area coverage and mission duration
Useful flight endurance	Maximum mission time before reserve/return
Altitude	Ground sampling distance and field coverage
Camera resolution / field of view	Spatial resolution and swath width
Image interval and overlap	Photogrammetry / coverage requirements

Engineering input	Use in simulation
Battery swap and setup time	Operational labor and downtime
Acres per flight	Capacity constraint
Processing time	Latency from flight to actionable information

5.1 Resolution-versus-coverage experiment

Simulate multiple altitudes or effective resolutions. Higher altitude should increase acres/hour but reduce spatial detail. The goal is not to maximize image resolution; it is to maximize the economic value of information per unit of flight effort.

Optimal operating point = $\text{argmax}(\text{economic information value})$

6. Phase 4 - Simulate RGB/NIR Observations Realistically

The cameras measure reflected light, not yield, disease, soil moisture, or nitrogen directly. The observation model should therefore generate indirect, noisy signals and avoid giving the drone perfect knowledge.

6.1 RGB-derived features

- Canopy coverage and vegetation segmentation
- Stand count / plant density where image resolution permits
- Row and emergence-gap detection
- Visible discoloration or canopy damage
- Texture / spatial pattern features
- Photogrammetric canopy height in later versions

6.2 RGB + NIR features

The primary vegetation index available from the current payload is NDVI:

$$\text{NDVI} = (\text{NIR} - \text{Red}) / (\text{NIR} + \text{Red})$$

NDVI should be one feature among several rather than the sole basis for crop-state estimation, particularly because it becomes less sensitive in dense vegetation. Combine spectral and structural features where possible.

6.3 Observation noise

A defensible virtual sensor should include error terms for:

- illumination variability and shadows
- sensor noise and exposure differences

- viewing-angle effects
- RGB/NIR co-registration error
- motion blur or missing images
- classification / segmentation error
- processing latency and occasional mission failure

7. Phase 5 - Convert Drone Observations into Crop-State Estimates

A practical technical bridge is to estimate LAI from RGB/NIR features, then use the estimated LAI to update WOFOST. Existing UAV/WOFOST research has used this general architecture with regression or machine-learning models and data assimilation.

RGB + NIR -> vegetation indices + canopy features -> regression / ML -> estimated LAI -> WOFOST / PCSE

Begin with interpretable, data-efficient models rather than large deep networks:

Model	Why include it
Linear / ridge regression	Transparent baseline and useful for calibration
Random Forest	Strong nonlinear baseline; robust with mixed features
XGBoost	Strong tabular predictor and useful performance benchmark
Neural network / CNN	Later phase if sufficient real imagery and labels exist

8. Phase 6 - Simulate Conventional Scouting Fairly

The human baseline should be explicit, reproducible, and strong. A scout should be modeled as highly accurate at a location they actually inspect, but spatially sparse and labor intensive.

Parameter	Illustrative implementation
Scouting interval	Every 7 days; later test other frequencies
Path	W-, Z-, or diamond-style path
Stops	Fixed number of locations per field
Plants inspected	Approximately 50-100 plants total per field in the baseline scenario
Local diagnostic accuracy	High when the scout physically reaches the affected location
Coverage	Sparse relative to field area

Parameter	Illustrative implementation
Travel / labor	Explicit walking, driving and recording time

Fair comparison

The drone does not have to outperform a human at close-range diagnosis. Its advantage can be whole-field coverage, repeatability, monitoring frequency, and the ability to direct a scout toward the locations most worth inspecting.

9. Phase 7 - Quantify Detection Performance

Every simulated stress event should have a known footprint. The conventional scout detects it only when the sampled path intersects the event and the local symptoms are observable. The drone detects it only when the area is imaged and the sensor/algorithm signal exceeds a threshold.

Metric	Definition
Detection rate	Detected stress events / total stress events
False-positive rate	Healthy zones incorrectly flagged / healthy zones
Time-to-detection	Detection date - true stress onset date
Minimum detectable anomaly size	Smallest stress footprint detected at the required probability
Localization error	Distance / area mismatch between true and estimated anomaly footprint

10. Phase 8 - Assimilate Drone Data into WOFOST/PCSE

Compare an open-loop crop model against a drone-updated model. The initial assimilation method should be an Ensemble Kalman Filter (EnKF) because it is widely used for sequential crop-state updating and is represented in PCSE example notebooks.

Model	Information flow
Open-loop WOFOST	Weather + soil + crop parameters + management -> predicted crop state / yield
Drone-assisted WOFOST	Open-loop inputs + repeated drone-estimated LAI -> EnKF update -> corrected crop state / yield forecast

Model	Information flow
Scout-assisted WOFOST	Open-loop inputs + sparse local crop observations -> update -> corrected forecast

11. Experimental Treatments

Use multiple information treatments so the contribution of each capability can be isolated.

Treatment	Information available	Purpose
T0 - No monitoring	Weather + initial farm assumptions	Baseline model uncertainty
T1 - Conventional scouting	Sparse field samples	Current-practice baseline
T2 - RGB drone	Canopy / visible anomaly features	Value of aerial coverage without NIR
T3 - RGB + NIR drone	RGB features + NDVI / NIR information	Incremental value of NIR
T4 - Drone + WOFOST assimilation	Drone observations update crop-state model	Value of data-model fusion
T5 - Perfect information	True simulated crop state	Upper bound on value of information

The perfect-information treatment is especially useful because it reveals what fraction of the theoretical maximum information value the drone captures.

12. Phase 9 - Convert Information into Decisions

Prediction accuracy alone is not customer value. A monitoring system becomes economically relevant when it changes an action. Add explicit decision policies only after the observation and estimation pipeline is working.

Decision case	Example policy
Targeted scouting	If anomaly score exceeds threshold, dispatch a scout only to flagged zones.
Irrigation	If predicted water stress exceeds threshold and avoided yield loss exceeds irrigation cost, irrigate.
Nitrogen	If inferred N stress exceeds threshold and growth stage

Decision case	Example policy
	permits response, apply targeted N.
No action	If predicted economic benefit is below intervention cost, continue monitoring.

12.1 Recommended first commercial workflow

The most defensible near-term workflow is drone-assisted targeted scouting rather than full autonomous diagnosis:

Whole-field drone survey -> anomaly map -> ranked inspection zones -> human diagnosis -> treatment decision

13. Phase 10 - Add the Economic Layer

Track economic value at the field and per-acre level. The whitepaper should emphasize net value created, not imagery quality alone.

Economic component	Example contents
Monitoring cost	Labor, travel, equipment allocation, batteries, processing
Input cost	Nitrogen, water, crop protection, application cost
Revenue	Harvested yield x crop price
Avoided loss	Yield preserved because an issue was detected and treated earlier
Net profit	Revenue - monitoring - inputs - other scenario-specific costs
Drone ROI	$(\text{Profit_drone} - \text{Profit_baseline}) / \text{drone monitoring cost}$

Sensitivity analysis should vary crop price, fertilizer price, labor cost, farm size, monitoring frequency, and drone cost so the output can answer questions such as “At what farm size does the system pay for itself?”

14. Phase 11 - Run Monte Carlo and Robustness Experiments

Do not base conclusions on one synthetic season. Run hundreds or thousands of seasons across changing environmental, operational, and economic conditions.

Randomized factor	Examples
Weather	Historical weather years; dry, normal, wet seasons
Soil state	Initial moisture, fertility, spatial heterogeneity
Stress event	Location, onset date, severity, footprint, duration
Monitoring constraints	Wind / rain no-fly periods, sensor noise, mission failure
Scouting	Path, visit timing, observation error
Economics	Crop price, labor rate, input prices

Preferred reporting style

Report distributions and confidence intervals rather than a single “saved \$X/acre” result. Example structure: across N simulated seasons, drone-assisted monitoring increased expected net profit by X and produced positive value in Y% of seasons.

15. Secondary Experiments

15.1 Monitoring frequency

Compare daily, every 3 days, weekly, biweekly, and growth-stage-triggered flights. Estimate the marginal value of an additional flight and identify the frequency beyond which added monitoring produces little incremental value.

15.2 Resolution versus acreage

Compare several altitudes/resolutions and calculate the operating point that gives the best detection and economic performance per flight hour, rather than simply the smallest ground sampling distance.

15.3 Farm scale

Run the same workflow at different farm sizes (for example 80, 160, 320, 640, and 1,280 acres) to estimate where operational savings and avoided loss justify the system cost.

15.4 Stress footprint and severity

Vary anomaly size and severity to determine the minimum economically meaningful problem the system can detect reliably.

16. KPI Framework

Category	Primary KPIs
Crop-state accuracy	LAI RMSE, crop-state bias, uncertainty reduction
Forecast accuracy	Yield RMSE / MAE / R ² ; forecast lead time
Detection	Detection rate, false positives, localization error, time-to-detection
Operations	Acres/hour, missions/field, labor hours/season, processing latency
Management	Time-to-intervention, input quantity, treated acreage
Economics	\$/acre monitored, avoided loss, net profit, ROI, payback threshold
Information value	Performance versus perfect-information upper bound

17. Minimum Viable Simulation (MVP)

The first implementation should deliberately omit complex disease diagnosis, hyperspectral sensing, autonomous treatment, and large deep-learning models. Build the smallest experiment capable of proving or falsifying the core value proposition.

CORN FIELD

```

|
+-- heterogeneous soil / nitrogen / moisture
+-- weather
|
v


```

WOFOST / PCSE -> TRUE DAILY LAI + YIELD

```

|
+-----+
|         |
|         |
v         v
FIELD SCOUT   VIRTUAL LOCUST
spatial samples  RGB + NIR
|               |
|               | NDVI + canopy features
|               |
|               | RF / regression
|               |
|               | estimated LAI
|               |
+-----+-----+
v

```



MVP research question

Testable question
Does crop information derived from the RGB/NIR drone improve field-level crop-state and yield estimation compared with conventional scouting at equal or lower monitoring labor?

18. Suggested Software Architecture

A modular Python design will make it easier to swap crop models, sensing assumptions, AI models, and management policies without rewriting the experiment.

Component	Responsibility
Farm / FieldZone	Spatial layout, soils, management zones, metadata
TruthModel	Runs the crop/soil/weather system that generates hidden true states
StressGenerator	Introduces controlled stress events
DroneMissionModel	Flight path, coverage, endurance, altitude, timing, no-fly constraints
DroneSensor	Converts true crop state to noisy RGB/NIR-derived observations
HumanScout	Samples field locations and produces local observations
FeatureModel	Computes synthetic VIs/canopy features
CropStateEstimator	WOFOST/PCSE estimator
Assimilator	EnKF or later alternative assimilation method
DecisionPolicy	Targeted scouting / irrigation / nutrient action rules
EconomicsModel	Costs, revenue, avoided loss, ROI

Component	Responsibility
ExperimentRunner	Scenario generation, Monte Carlo runs, metrics, reproducibility

19. Implementation Roadmap

Stage	Deliverable	Exit criterion
1. Spatial crop model	Multi-zone WOFOST/PCSE field	Different zones produce plausible crop trajectories
2. Stress injection	Known spatial water/N/stand-loss events	True onset, footprint and yield effect are measurable
3. Human scout model	Repeatable sparse sampling baseline	Detection/labor metrics are generated
4. Drone mission model	Coverage + endurance + timing based on real aircraft data	Acres/hour and mission constraints match physical tests
5. Virtual RGB/NIR sensor	Noisy NDVI / canopy features	Sensor outputs are plausible and tunable
6. LAI estimator	Regression/RF/XGBoost benchmark	Cross-validated LAI error established
7. EnKF assimilation	Drone-updated WOFOST state	Open-loop vs assisted forecast can be compared
8. Monte Carlo	Large scenario batch	Performance distributions are stable
9. Decision/economic layer	Targeted scouting / treatment policies	Net value and ROI can be calculated
10. Model-mismatch validation	Truth differs from estimator	Final results are not circular

20. Whitepaper Claims This Framework Can Support

If the experiment is implemented carefully, it can support bounded, evidence-based claims such as:

- How much faster the drone detects spatial anomalies under specified monitoring frequencies.
- How monitoring coverage and labor compare with conventional scouting.
- How much RGB/NIR observations reduce uncertainty in crop-state or yield estimates.

- Whether NIR adds measurable value beyond RGB for the selected crop/stress cases.
- How effectively drone surveys direct a human scout toward high-value inspection zones.
- How flight frequency and altitude affect the economic value of information.
- The conditions under which drone-assisted monitoring produces positive net value or ROI.

The framework should not claim that RGB/NIR imagery directly diagnoses every disease, pest, nutrient deficiency, or soil condition unless that capability is separately validated against real labeled field data.

21. Methodological Risks and Controls

Risk	Control
Circular validation	Use model mismatch or a different truth model for final experiments
Perfect virtual sensor	Inject realistic measurement, classification and operational error
Weak conventional baseline	Give the human scout high local diagnostic accuracy and explicit sampling logic
Overfitting AI	Use held-out data / cross-validation and simple baseline models
Unrealistic management value	Require an explicit action pathway before assigning avoided loss
Single-season conclusions	Use Monte Carlo / multiple weather years
Overclaiming causality	Report assumptions and conduct sensitivity analysis

22. Recommended Technical References

The following resources are directly relevant to building and validating this simulation architecture:

- **PCSE repository and documentation**
<https://github.com/ajwdewit/pcse>
- **PCSE example notebooks (including assimilation / calibration examples)**
https://github.com/ajwdewit/pcse_notebooks
- **Improving maize yield estimation by assimilating UAV-based LAI into WOFOST (2024)**
<https://www.sciencedirect.com/science/article/pii/S0378429024002302>
- **Maize yield estimation using UAV multispectral monitoring and WOFOST data assimilation (2025)**
<https://www.sciencedirect.com/science/article/abs/pii/S1161030125001108>
- **Systematic review of remote-sensing data assimilation into crop models**
<https://www.sciencedirect.com/science/article/abs/pii/S0308521X23001166>

- **NASA - vegetation indices / NDVI background**
<https://science.nasa.gov/earth/earth-observatory/measuring-vegetation-ndvi-evi/>
- **Iowa State Extension - crop scouting basics**
<https://crops.extension.iastate.edu/encyclopedia/crop-scouting-basics>
- **Iowa State Extension - aerial imagery for crop scouting**
<https://crops.extension.iastate.edu/post/crop-scouting-aerial-imagery-choosing-right-source>

23. Final Recommended Research Statement

Proposed whitepaper thesis

High-frequency RGB/NIR UAV observations can reduce uncertainty in crop-state estimation relative to conventional spatial sampling. The business value of the drone should be evaluated by how that information changes anomaly detection, targeted scouting, management timing, input use, yield loss, and net farm economics - not by image quality alone.

A successful MVP will establish the information advantage. A successful Version 2 will establish the decision advantage. A successful Version 3 will establish the predictive and economic advantage of maintaining a repeatedly updated digital representation of crop state across the field.